



**Batch: C-1      Roll No: 16010121031**

**Experiment / assignment / tutorial No: 5**

**Grade: AA / AB / BB / BC / CC / CD / DD**

**Signature of the Staff In-charge with date**

**TITLE:** To implement temporal difference approach

**AIM:** TD learning and prediction

**Expected OUTCOME of Experiment:**

CO-3: Apply different temporal difference learning policies.

**Books/ Journals/ Websites referred:**

Richard S. Sutton and Andrew G. Barto, "*Reinforcement Learning: An Introduction*", The MIT Press, Second Edition, 2018

**Pre-Lab/ Prior Concepts:**

TD methods learn directly from episodes of experience. TD is model-free: no knowledge of MDP transitions / rewards is required. TD learns from incomplete episodes, by bootstrapping. TD updates a guess towards a guess. TD can learn before knowing the final outcome. TD can learn online after every step.

**Chosen Problem Statement:** Driving Home

**Explain following concepts w.r.t. chosen problem statement:**

▪ **Policy:**

The policy defines how the driver (agent) selects actions at each decision point, such as choosing routes or adjusting speed. In an epsilon-greedy policy:

- Exploitation* ( $1 - \epsilon$ ): The driver chooses the route that minimizes expected travel time based on current knowledge.
- Exploration* ( $\epsilon$ ): The driver explores alternative routes to potentially discover faster or more efficient paths.

▪ **Reward Function:**

The reward function provides feedback based on travel time or cost:

- Negative Reward:* The agent receives a negative reward proportional to the time taken for each route segment. Longer routes incur larger penalties (negative rewards).
- Positive Reward:* Reaching home within a specific time frame may yield a positive reward, incentivizing the agent to optimize driving time.

- **Value** **Function:**  
The value function estimates the expected total reward, typically minimizing travel time, starting from a specific state:
  - a) *State Value Function* ( $V(s)$ ): The expected remaining travel time starting from state  $s$ .
  - b) *Action Value Function* ( $Q(s, a)$ ): The expected remaining travel time when taking action  $a$  in state  $s$ .
- **Model of the Environment:**  
The environment consists of external factors that influence driving decisions, such as:
  - a) *Road Network*: The routes available between the starting point and home.
  - b) *Traffic Conditions*: Varying traffic that can delay or speed up travel on different routes.
  - c) *Weather*: External conditions like rain or snow that may affect driving speeds.
  - d) *Signals and Rules*: Traffic lights, stop signs, and speed limits that impact the time taken at intersections.

### Implementation

#### Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Define states
states = ['A', 'B', 'C', 'D', 'E', 'F']
num_states = len(states)

# Define initial elapsed times with a 5 min gap
initial_elapsed_times = [0, 5, 10, 20, 30, 40]

# Add noise to the elapsed times
np.random.seed(42)
noise_elapsed = np.random.normal(0, 1, num_states)
final_elapsed_times = np.maximum(np.array(initial_elapsed_times) +
                                  noise_elapsed, 0)

# Define fixed predicted times from each state to the final state
predicted_times_to_go = [30, 25, 20, 15, 10, 0]

# Add noise to the predicted times
noise_predicted = np.random.normal(0, 2, num_states)
final_predicted_times_to_go =
np.maximum(np.array(predicted_times_to_go) + noise_predicted, 0)
```

```
# Calculate total predicted times
total_predicted_times = final_elapsed_times +
final_predicted_times_to_go

# Round the times to two decimal places
final_elapsed_times = np.round(final_elapsed_times, 2)
final_predicted_times_to_go = np.round(final_predicted_times_to_go,
2)
total_predicted_times = np.round(total_predicted_times, 2)

# Initialize value function
V = {state: 0 for state in states}

# Parameters
alpha = 0.1 # Learning rate
gamma = 1.0 # Discount factor
num_episodes = 1000

# TD(0) Learning
for episode in range(num_episodes):
    state_idx = 0
    while state_idx < num_states - 1:
        state = states[state_idx]
        next_state = states[state_idx + 1]
        reward = -total_predicted_times[state_idx] # Negative to
minimize time
        if next_state == 'F':
            V[state] += alpha * (reward - V[state])
        else:
            V[state] += alpha * (reward + gamma * V[next_state] -
V[state])
        state_idx += 1

# Create a DataFrame for tabular output
data = {
    'State': states,
    'Elapsed Time (minutes)': final_elapsed_times,
    'Predicted Time to Go (minutes)': final_predicted_times_to_go,
    'Predicted Total Time (minutes)': total_predicted_times
}
df = pd.DataFrame(data)

# Print the DataFrame
print(df)

# Plot the graph
plt.figure(figsize=(10, 6))
plt.plot(states, total_predicted_times, marker='o')
```

```

plt.xlabel('State')
plt.ylabel('Total Predicted Time (minutes)')
plt.title('Total Predicted Time vs. State using TD(0) RL')
plt.grid(True)
plt.show()

```

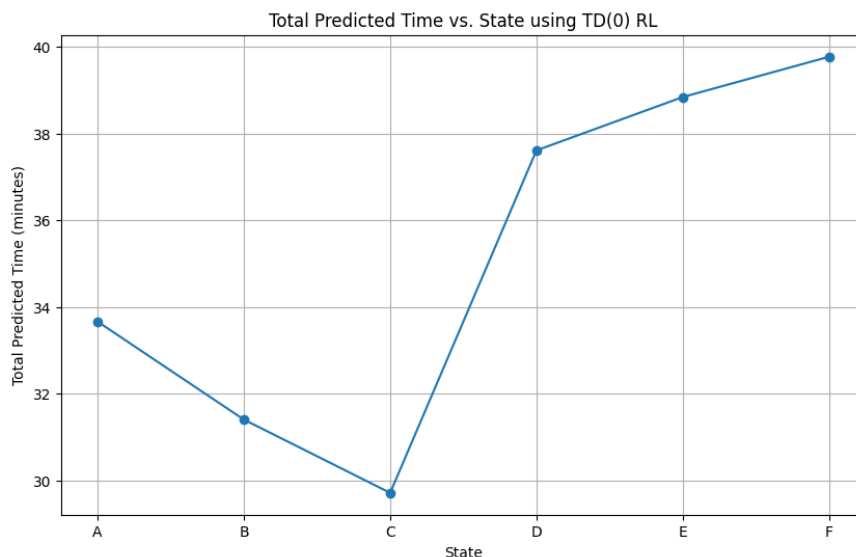
**Output:**

```

State Elapsed Time (minutes) Predicted Time to Go (minutes) \
0      A                   0.50                      33.16
1      B                   4.86                      26.53
2      C                  10.65                      19.06
3      D                  21.52                      16.09
4      E                  29.77                       9.07
5      F                  39.77                       0.00

Predicted Total Time (minutes)
0                      33.66
1                      31.40
2                      29.71
3                      37.61
4                      38.84
5                      39.77

```



### **Conclusion:**

Successfully implemented the Driving Home problem using Temporal difference learning.

### **Post Lab Descriptive Questions:**

**Q - List advantages and disadvantages of MC and TD approach.**

Aspect	Monte Carlo (MC)	Temporal Difference (TD)
<b>Learning Style</b>	Learn from complete episodes (episodic).	Learn from each step (online/bootstrapping).
<b>Exploration</b>	Requires exploration across episodes.	Learns with each step, can adapt faster to exploration changes.
<b>Convergence</b>	Converges slowly, only after complete episodes.	Typically converges faster since updates happen more frequently.

<b>Computation</b>	Requires the entire episode to be stored and computed, which can be memory-intensive.	Updates incrementally, lower memory requirements.
<b>Variance</b>	High variance due to reliance on complete episodes.	Lower variance as it updates after each step.
<b>Bias</b>	Unbiased estimates of the return since it's based on full episodes.	Can introduce bias due to bootstrapping from estimates.
<b>Applicability</b>	Only works for episodic tasks.	Works for both episodic and continuous tasks.
<b>Handling of Non-Markovian Environments</b>	Performs poorly in non-Markovian environments as it waits until the end of an episode.	Handles non-Markovian environments better with step-by-step updates.

**Date:**

**Signature of faculty in-charge**